

Render · Career Launch Dashboard

A student-facing AI career readiness prototype for Digital Media Arts graduates

Version: 0.5 · Status: **Prototype, preparing for Fall 2026 pilot** · May 15, 2026

Michelle Blomberg · Digital Media Arts Faculty · Glendale Community College

1. Executive Summary

The Problem

Creative students at community colleges graduate with strong portfolios but no systematic plan for using them. They lose Canvas access after graduation, and with it, much of their coursework. There is no career-launch infrastructure in place. A typical pattern: students graduate, ask “now what?”, and lose momentum at the most critical moment of their career launch.

The Intended Solution

Render is a student-facing AI career dashboard intended to be populated during AVC 248 (Design Self Promotion at GCC) and owned by the student after graduation. Students don't face a blank AI prompt. They interact with a structured tool designed for their specific situation, building a career launch system they actually know how to use.

Current state. Render is a working prototype. All seven planned phases have been coded. The tool has not yet been used by a full class. Small-scale usability feedback was gathered in Spring 2026. First real classroom deployment is planned for Fall 2026 in AVC 248.

Intended outcome at graduation. Students leave with a populated career launchpad: goals, job search strings, links, resume drafts, job search history with cover letters, skills gap analysis, networking contacts, interview practice records, a weekly schedule, and a 90-day professional development plan. All exported as a standalone HTML file they own. Whether this works as intended is the question the Fall 2026 pilot is designed to answer.

2. Product Vision

2.1 Vision Statement

What if students left your course not just with completed assignments, but with a personalized career launch system they built themselves, and know how to keep using?

Render is the intended answer to “now what do I do?”. A career dashboard students build incrementally across a semester and take with them when they graduate. The post-graduation dashboard is designed to be a living document, not a static export, that students can continue to update as their career develops. The post-graduation usage pattern is a design goal. It has not been validated by graduates yet.

2.2 Two-Product Architecture

#	Product	Purpose	Status
1	In-Class Gathering Interface	Structured data entry tool intended for use during AVC 248. Students populate it week by week as assignments are completed. AI assists with analysis, cover letters, search strings, and coaching.	<i>Prototype coded</i>
2	Post-Graduation Dashboard	A persistent, updatable career dashboard exported from Product 1. The intent is for students to continue using and updating it after graduation. 4 color themes. No login required. Includes weekly schedule and PD plan.	<i>Prototype coded</i>

3. Users & Context

3.1 Primary Users

User	Context	Needs from Render
DMA students	Finishing AVC 248. Portfolios done. May be seeking employment, freelance clients, or both.	Structured guidance, not blank prompts. Somewhere to put work they've already done. A plan for what comes next.
Animation students	Portfolio-focused. Targeting studios, production companies, streaming platforms. Reel is primary job search tool.	Reel link prominent. Industry-specific job analysis. Animation software skill tracking.
Photography students	Mix of commercial and fine art. Client-finding as relevant as job-searching. Instagram and portfolio site are key.	Freelance client track. Social media links. Commercial vs. editorial distinction.
Video production	Targeting agencies, corporate video, broadcast, or building a freelance production company.	Both employment and freelance tracks. Reel + portfolio. Spec sheet and contract templates.
UX design students	Most likely pursuing employment at tech companies. Case study portfolio is critical. LinkedIn is primary platform.	Job search focus. Case study links. UX-specific skill gap analysis.
Instructor (Michelle B.)	Runs AVC 248. ~24 students per semester. Interested in sharing employer signal with Career Services if students opt in.	Instructor view of class usage. Configurable data pipeline. Privacy-safe defaults.
Career Services (GCC)	In conversation about whether anonymous employer interest data from Render could help build internship pipelines. Mollie is the contact. Pipeline is proposed, not formally adopted.	If they choose to use it: anonymous aggregated data, employer names, job titles, programs. No student names or MEIDs.

3.2 Course Context

- Course: AVC 248 Design Self Promotion. 15 weeks, online, asynchronous.
- Class size: approximately 24 students per semester.
- Students come in with portfolios already done. They need to know how to use them.
- Canvas access ends after graduation. Render must be self-contained once exported.
- Student devices: mix of Mac, Windows, and mobile. Tool must be device-agnostic.
- First classroom use is planned for Fall 2026. Spring 2026 was small-scale usability feedback only.

4. Features & Build Phases

All seven phases are coded in the current prototype. None have been validated with a full class. The Fall 2026 pilot will determine which work as intended, which need rework, and which the students don't actually use.

Phase 1, Profile + Goals + Links Hub + Search Strings (Weeks 1–3)

- First name only. No email, no MEID, no personally identifiable information collected.
- Program selection, employment/freelance track, weekly commitment sliders, links hub (LinkedIn, Behance, portfolio, reel, Instagram, GitHub).
- Goals & identity statements, dream job/client, target market, 3-year vision.
- AI Search String Generator. Generates ready-to-paste strings for LinkedIn, Indeed, Glassdoor, and Google site: searches. Each string has an individual copy button.

Status: *Coded, awaiting classroom validation in Fall 2026.*

Phase 2, Job / Client Log (Weeks 2–13)

- Employment tab: paste job description, AI skills gap analysis, AI cover letter generation, status tracker (Saved → Applied → Interview → Offer/Pass). Full job description, resume edits, and cover letter stored together per entry.
- Freelance tab: describe ideal client type, AI generates search strategies and 8–10 prospect names, outreach guidance, pricing check.
- Proposed Career Services pipeline: when a student saves a job entry, the prototype sends employer name, job title, and program area to a shared Google Sheet. Students see this transparency notice before saving. No student names or personal information are included. Whether Career Services formally adopts this pipeline is an open question.

Status: *Coded. Pipeline mechanics work end-to-end in testing. Adoption by Career Services is pending.*

Phase 3, Resume Vault (Weeks 4–9)

- Base resume drafts (1st through 4th) with instructor feedback fields and active resume flag. Each job entry holds tailored AI resume edit suggestions alongside the cover letter.
- Privacy note displayed to students: resume drafts in Render should not include name, address, phone number, or other personal contact information. When a student is ready to send a resume to an employer, they add contact details to the version they send, not the version in Render.

Status: *Coded, awaiting classroom validation in Fall 2026.*

Phase 4, Skills & PD (Weeks 5–13)

- Software stack checklist by program (6 categories, approximately 60 tools).
- Skills-to-build list. Add manually or pull gaps from saved job entries via AI. Gap analysis surfaces professional skills and industry workflows (project management, client presentations, creative briefs, production planning, budget management) alongside software tools.
- Named free learning resources, including AI-skills resources alongside the traditional software tutorials. Specific YouTube channels, Coursera courses, Adobe tutorials, etc.
- Portfolio project idea generator. PD activity log.

Status: *Coded. Resource quality needs validation in Fall pilot.*

Phase 5, Networking Tracker (Weeks 11–12)

- Contact log with type tagging (mentor, employer, peer, alumni, industry).
- Outreach status tracker.
- AI drafts personalized LinkedIn connection requests under 300 characters. Redraft capability per contact.

Status: *Coded, awaiting classroom validation in Fall 2026.*

Phase 6, Interview Prep (Week 13)

- Role-specific questions generated from saved job descriptions.
- General questions by role type.
- Answer practice with AI coaching feedback. Coach, not cheerleader.
- Saved Q&A log.

Note on Big Interview: *Students currently have access to Big Interview through GCC during the class. Big Interview access ends when the class does. Render's Interview Prep is not integrated with Big Interview, and it is not intended to be. Render is a separate practice option that students keep after they lose Big Interview access.*

Status: *Coded as text-only practice. See §11.1.1 for the Summer 2026 v1 plan to add voice mode.*

Phase 7, Launch Plan (Week 15)

- AI 90-day professional development plan generator based on accumulated student data.
- Semester stats at a glance.
- AI Weekly Schedule Builder. Generates a Monday–Sunday time-blocked schedule based on the student's committed hours, goals, skill gaps, and job targets. Includes specific activities, platforms, networking touches, and a Sunday review block.
- Product 2 export with 4 color themes (Floral, Studio Dark, Clean Minimal, Bold Editorial). Export includes search strings and weekly schedule.

Status: *Coded, awaiting classroom validation in Fall 2026.*

5. Data Architecture & Privacy

5.1 Student Privacy Commitment

Render collects only a first name. No email, no student ID, no address, no phone number. Students are understandably skeptical about AI tools and data collection. The prototype is designed to earn trust by collecting the absolute minimum and being transparent about what happens with every piece of data.

Resume drafts stored in Render should not include personal contact information. A visible note reminds students: “When you’re ready to send this to an employer, add your name and contact info to the version you send, but keep them out of Render.”

Job search entries trigger a transparent notification. Students are told that employer name, job title, and program area would be shared with Career Services if the pipeline is adopted. No student-identifying information is included.

5.2 Data Storage

All student data in the current prototype is stored in browser localStorage. Data stays on the student’s own device. This is intentional for the prototype and the Fall 2026 pilot.

- No server, no database, no third-party data storage in current build.
- Students export their data as a JSON file saved to their own Google Drive for backup.
- Auto-save runs every 60 seconds within the session.
- Future migration path under consideration: Maricopa Google infrastructure or school server. Not built. Not committed.

5.3 FERPA Considerations

Data type	Where it goes	FERPA status
First name only, goals, job history, resume text without personal contact info	Browser localStorage only. Stays on student’s device.	<i>Safe. Never transmitted.</i>
Employer name, job title, URL, interest level, program area	Google Sheet via Apps Script (proposed Career Services pipeline)	<i>Safe. No student identifiers transmitted.</i>
Full student data export (JSON)	Student’s own Google Drive. Student controls it.	<i>Safe. Student owns the data.</i>
Future: cloud sync (not built)	Maricopa Google infrastructure or school server (under consideration)	<i>Would require FERPA agreement. Not yet pursued.</i>

5.4 Career Services Pipeline (Proposed)

When a student saves an employment job entry, the prototype currently sends an anonymous row to a Google Sheet via Google Apps Script. The intent is for Career Services to use this data to identify employers students are interested in, build relationships with those companies, and pursue internship and placement opportunities.

This pipeline works technically. Whether Career Services adopts it as an active workflow is an open conversation with Mollie in Career Services. Until adoption is confirmed, the pipeline is considered proposed, not operational.

Data fields sent (if adopted): Date, Employer Name, Job Title, Job URL, Interest Level, Program Area.

Not sent: student name, first name, MEID, or any personally identifiable information.

6. Technical Specification

6.1 Current Stack

- Frontend: single HTML file. Vanilla JavaScript, CSS, no framework.
- Authentication: first name only (prototype). No student ID, no email, no password.
- Persistence: browser localStorage, keyed by first name.
- AI: Anthropic Claude API (claude-sonnet-4-20250514) via direct fetch calls.
- Career Services pipeline: Google Apps Script Web App (POST endpoint, no-cors mode).
- Hosting: GitHub Pages (singletrackmom.github.io/render/).
- Export: JSON file download. Google Drive auto-sync is under consideration, not built.

6.2 AI Functions

Fifteen AI functions are wired into the prototype. All calls include student profile and goals as context. Model: claude-sonnet-4-20250514.

AI Function	Input	Output
analyzeWithAI()	Profile + goals + job description	Skills match, gaps, positioning advice, red flags
generateCoverLetter()	Profile + job desc + analysis	Tailored 3-paragraph cover letter
sendToSheet()	Job description (up to 1500 chars)	2-3 sentence summary to proposed Career Services sheet
getResumeEdits()	Active resume + job description	Specific resume edits for this job
pullGapsFromJobs()	Saved job descriptions (up to 6)	8-12 skills: software, professional, workflows
getRecommendations()	Profile + skills gap list	Named free learning resources
getProjectIdeas()	Profile + targets + gaps	5 portfolio project ideas with time estimates
draftLinkedInMessage()	Contact info + notes	Personalized connection request under 300 chars
redraftMessage()	Contact + previous message	Fresh message for existing contact
generateInterviewQuestions()	Profile + job description	10 role-specific questions

generateGeneralQuestions()	Profile + role type	8 questions for any creative role
getFeedback()	Question + student answer	What's working, what to improve, stronger version
generatePDPlan()	All student data + hours + focus	Personalized 90-day plan
generateSearchStrings()	Goals + dream job + market	Copy-paste strings for LinkedIn, Indeed, Glassdoor, Google
buildWeeklySchedule()	Hours + goals + gaps + targets	Mon-Sun time-blocked schedule

7. Product 2, The Render Dashboard

7.1 Overview

Product 2 is the persistent career dashboard students keep and are intended to continue using after graduation. It is generated from data accumulated in Product 1. Students lose Canvas access after graduation, so Product 2 is what they own and control going forward.

Whether students actually keep using Product 2 post-graduation is a design assumption that needs validation. The Fall 2026 pilot will produce a cohort of graduates whose post-graduation usage can be observed.

Contents: goals, search strings, links hub, job search log with cover letters and resume edits, active resume, skills map with gaps and learning resources, PD activity log, network contacts, interview Q&As with AI feedback, 90-day PD plan, weekly schedule, and semester stats.

7.2 Theme Options

Product 2 includes four customizable color and typography palette options.

#	Theme	Colors	Typography
1	Floral (default)	Plum, sage, rose, gold on warm paper. Matches the AVC 248 site.	Lora + DM Sans
2	Studio Dark	Near-black, warm white, electric accent. Editorial / agency feel.	Playfair Display + DM Sans
3	Clean Minimal	White background, charcoal type, single accent. Corporate / tech.	Inter + Inter
4	Bold Editorial	High contrast, large type, strong blocks. Motion / entertainment.	Bebas Neue + DM Sans

Sample dashboard with realistic fake data: singletrackmom.github.io/render/sample-dashboard.html

8. Testing Status

Honest current state, May 2026:

- All seven phases are coded in the prototype.
- The prototype has not been used by a full AVC 248 class yet.
- Small-scale usability feedback was gathered informally in Spring 2026. Not a structured study, not a representative sample, not a substitute for a real pilot.
- Career Services conversation with Mollie is ongoing. Pipeline is technically functional, adoption is not yet confirmed.
- First real classroom deployment: Fall 2026 in AVC 248.
- Summer 2026 is for iteration toward v1: bug fixes from Spring usability feedback, three planned new features (see §11), and end-to-end testing with three simulated personas before students touch it.

9. Open Questions

Question	Priority	Notes
How should authentication work long-term? First-name-only is ideal for privacy but limits personalization across sessions.	High	Personal email + password on a real server is the likely path post-pilot. The no-institutional-ID principle should be preserved.
Where does student data live after pilot? localStorage is device-specific and not durable.	High	Server-side storage needed eventually. Current workaround: JSON export/import via Google Drive.
Will graduates actually keep using Product 2 after they leave?	High	This is the core design assumption. The Fall 2026 cohort is the first chance to observe post-graduation usage.
Should the tool be shared across MCCCD or remain GCC-specific?	Medium	Architecture is abstracted. District-wide deployment is conceptually viable. Active conversation with the District AI CIO. See §11.1.3.
Will Career Services actively use the Google Sheet pipeline?	Medium	Pending confirmation with Mollie. Pipeline is built but unused.
Should there be an instructor view showing class-wide progress?	Low	Useful for grading. Would require explicit student consent for FERPA compliance.

10. Roadmap

Milestone	Target	Notes
All seven phases coded in prototype	March 2026	Complete. Coded, not validated with a real class.

Small-scale usability feedback (Spring 2026)	March–April 2026	Complete. Informal, not a structured study.
Summer 2026 iteration toward v1	June–Aug 2026	Bug fixes, three new features (see §11), personas-based end-to-end testing.
Pilot, AVC 248	Fall 2026	First full-class deployment. Feedback collection and iteration.
Google Drive auto-sync	<i>Under consideration</i>	Cross-device data transfer without manual export/import.
Canvas LTI integration	<i>Under consideration</i>	Embed Render inside Canvas rather than as a separate link.
MCCCD district-wide adaptation	<i>Under consideration</i>	Deployment beyond GCC if Fall 2026 pilot succeeds. See §11.1.3.

11. Upcoming Features (Coming Soon)

11.1 Summer 2026 v1 Build Cycle

Three features are planned for the Summer 2026 build cycle. All three were scoped during the Spring 2026 usability feedback period. These are planned, not yet started.

11.1.1 Voice-mode Interview Coach

The current Interview Coach is text-based. Students type their answers and the LLM provides STAR-method coaching feedback. The voice-mode build would add a spoken-response layer using the Web Speech API for browser-native transcription, plus client-side audio analysis to flag filler words (“um”, “like”, “you know”) and pacing patterns.

Intended output: per-answer feedback that grades content, delivery (filler density, pause patterns, words per minute), and structure (presence of Situation/Task/Action/Result components).

On Big Interview: students currently have Big Interview access through GCC during the class only. Render’s voice-mode Coach is not intended to compete with Big Interview during the class. It is intended to be a practice option students still have after the class ends, when Big Interview access goes away.

Technical scope: Web Speech API for browser-side speech-to-text, lightweight client-side audio analysis (no server upload), prompt re-engineering for the multimodal context. Estimated build time: 3 weeks.

11.1.2 Live Company Research Agent

An agentic enhancement to the Interview Coach. When a session starts on a saved job, the agent would use the LLM’s tool-use capability to audit the target company’s actual website, recent news, and public social presence in real time. The retrieved context would shape the question set and grade the student’s responses to the closing “Do you have any questions for me?” on whether they demonstrated genuine company research.

This would be a step up in agentic complexity from the current build. The Interview Coach today is a deterministic multi-turn loop with full context known up front. The company research agent would add dynamic information retrieval (tool use), multi-source synthesis, and conditional question generation

based on what the agent learns about the company. This would be the first feature in Render that justifies the term “agent” in its strict technical sense: autonomous tool selection within a loop.

Technical scope: Anthropic Claude tool-use API, web fetch tool, summarization prompt chain, prompt updates throughout the Interview Coach module. Estimated build time: 4 weeks.

11.1.3 District-wide Maricopa Rollout

Adapting Render for use across all ten Maricopa Community College District campuses. A conversation with the District AI CIO is in progress. The proposed architecture: school-picker pattern at launch (student selects their home college), program-specific course schedules and resource libraries, shared core tool. Critically, no central data store. Every student’s data still lives locally on their device.

Pedagogically, this requires that each college’s career services pipeline be configurable. Resource libraries (job boards, professional associations, industry-specific training) vary significantly by program. The architecture supports this through a JSON configuration file per program, not a code change. Programs could be added by faculty without engineering involvement.

Whether this happens depends on the Fall 2026 GCC pilot results. The district rollout is a possibility, not a commitment.

Technical scope: refactor program-specific data into external JSON, add school-picker UI, build resource library configuration schema, create documentation for adopting faculty. Estimated build time: 6 weeks.

11.2 Longer-Term Considerations

- Google Drive auto-sync for cross-device data transfer without manual export/import.
- Canvas LTI integration. Embed Render inside Canvas rather than as a separate link.
- Industry news feed. Curated RSS sources tailored to each student’s program track, so students follow industry news throughout the semester.
- Open-source release. Publish Render as a community-college EdTech reference implementation for other faculty to fork and adapt.

Appendix: Key References

- Render prototype: singletrackmom.github.io/render/
- Render overview (portfolio): singletrackmom.github.io/render/overview.html
- Sample exported dashboard: singletrackmom.github.io/render/sample-dashboard.html
- AVC 248 course site: singletrackmom.github.io/canvas/avc248/
- Big Interview (GCC, during-class access only): gccaz.biginterview.com
- Career Services: gccaz.edu/student-life/career-services